

Model Evaluation

Model evaluation is the process of measuring how well a machine learning model performs on unseen data. It ensures that the model is not just memorizing training data but can generalize new inputs.

It is important because it helps:

- Compare different models
- Detect problems like overfitting or underfitting
- Choose the best model for real-world use

For classification tasks, common evaluation tools include the confusion matrix, accuracy, precision, recall, and ROC-AUC.

Confusion Matrix

A confusion matrix is a table used to evaluate the performance of a classification model by comparing actual and predicted values. It gives a detailed breakdown of correct and incorrect predictions, which helps in understanding model errors.

In binary classification, it contains four outcomes:

- **True Positive (TP):** The model predicts positive, and the actual class is also positive.
- **True Negative (TN):** The model predicts negative, and the actual class is also negative.
- **False Positive (FP):** The model predicts positive, but the actual class is negative. (Type I error)
- **False Negative (FN):** The model predicts negative, but the actual class is positive. (Type II error)

Unlike accuracy, it shows **which type of mistakes the model is making**, making it very useful for analysis.

Example

Suppose a model predicts whether an email is spam or not spam. Out of 100 emails:

- 40 are spam
- 60 are not spam

The model correctly identifies:

- 32 spam emails $\rightarrow$ TP = 32
- 50 non-spam emails $\rightarrow$ TN = 50

So:

- FP = 10
- FN = 8

Confusion Matrix Table

Actual \ Predicted	Spam	Not Spam
Spam	32	8
Not Spam	10	50

Accuracy

Accuracy is the ratio of correct predictions to total predictions. It gives an overall idea of model performance.

$$\text{Accuracy} = (\text{TP} + \text{TN}) / \text{TP} + \text{TN} + \text{FP} + \text{FN}$$

Using the example:

$$\text{Accuracy} = (32 + 50) / 100 = 0.82 = 82\%$$

Explanation:

Accuracy is easy to understand but can be misleading when data is imbalanced. For

example, if 95% of emails are not spam, a model predicting all emails as "not spam" will still achieve 95% accuracy but fail in actual detection.

Precision

Precision measures how many predicted positive cases are actually correct. It focuses on the quality of positive predictions.

$$\text{Precision} = \text{TP} / (\text{FP} + \text{TP})$$

Using the example:

$$\text{Precision} = 32 / (32 + 10) = 32 / 42 \approx 0.76 = 76\%$$



Explanation:

High precision means fewer false positives. It is important in cases where wrong positive predictions are costly, such as:

- Spam filtering (important emails should not be marked spam)
- Fraud detection (false alarms should be minimized)

Recall

Recall measures how many actual positive cases are correctly identified by the model.

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$$

Using the example:

$$\text{Recall} = 32 / (32 + 8) = 0.80 = 80\%$$



Explanation:

High recall means fewer false negatives. It is critical in situations where missing a positive case is dangerous, such as:

- Disease detection
- Security threats

ROC Curve and AUC

The ROC (Receiver Operating Characteristic) curve is a graphical representation of a model's performance. It plots:

- **True Positive Rate (TPR)**
- **False Positive Rate (FPR)**

at different threshold values.

- $\text{TPR} = \text{TP} / (\text{TP} + \text{FN})$
- $\text{FPR} = \text{FP} / (\text{FP} + \text{TN})$

👉 It helps visualize the trade-off between sensitivity (recall) and false alarms.

AUC (Area Under Curve)

AUC summarizes the ROC curve into a single value.

- AUC close to **1** → Excellent model
- AUC around **0.5** → Random guessing
- AUC < **0.5** → Poor model

Explanation:

A higher AUC means the model can better distinguish between positive and negative classes regardless of threshold.

Example

In disease detection:

- Lower threshold → more patients detected (high recall) but more false positives
- Higher threshold → fewer false positives (high precision) but may miss actual cases

Unsupervised Learning

Unsupervised learning is a type of machine learning where the model works on unlabeled data. It automatically finds hidden patterns, structures, or relationships in the data.

It is useful when:

- Labels are not available
- We want to explore unknown data
- We need grouping or feature reduction

Main tasks include:

- Clustering
- Dimensionality reduction

Clustering

Clustering groups similar data points together based on their features. The goal is to maximize similarity within clusters and minimize similarity between clusters.

👉 It is widely used for pattern discovery and data exploration.

Applications

- Customer segmentation
- Image grouping
- Document clustering
- Anomaly detection

K-Means Clustering

K-Means is a popular clustering algorithm that divides data into **K clusters**, where K is predefined.

It works by minimizing the distance between data points and their cluster centroids.

Steps of K-Means

1. Choose number of clusters K
2. Initialize K centroids randomly
3. Assign each data point to nearest centroid
4. Update centroids by calculating mean
5. Repeat until convergence

👉 The goal is to minimize **within-cluster variation**.

Example

A shop groups customers based on spending and visits into:

- Low-value customers
- Regular customers
- High-value customers

Advantages

- Simple and easy to implement
- Fast for large datasets
- Works well when clusters are clear

Limitations

- Requires predefined K
- Sensitive to outliers
- Not suitable for irregular-shaped clusters

Hierarchical Clustering

Hierarchical clustering builds a tree-like structure (dendrogram) showing relationships between data points.

It does not require pre-defining the number of clusters.

Types

- **Agglomerative:** Starts with individual points and merges them
- **Divisive:** Starts with one cluster and splits it

Example

Countries grouped based on:

- GDP
- Literacy rate
- Life expectancy

Advantages

- No need to specify clusters initially
- Provides clear visualization

Limitations

- Computationally expensive
- Cannot undo merging/splitting

K-Means vs Hierarchical Clustering

Feature	K-Means	Hierarchical
---------	---------	--------------

Need beforehand	K	Yes	No
Output		Flat clusters	Tree (dendrogram)
Speed		Fast	Slow
Interpretability		Moderate	High

Dimensionality Reduction: PCA

Dimensionality reduction reduces the number of features while preserving important information. It helps simplify data and improve model performance.

Principal Component Analysis (PCA)

PCA transforms original features into new variables called **principal components**, which capture the maximum variance in the data.

Basic Steps

1. Standardize data
2. Compute covariance matrix
3. Find eigenvalues and eigenvectors
4. Select top components
5. Transform data

Example

An image dataset with 784 features can be reduced to 50–100 features using PCA, making processing faster and reducing noise.

Advantages

- Reduces computation time
- Helps visualize high-dimensional data
- Removes redundancy

Limitations

- Only captures linear relationships
- Hard to interpret new features
- Sensitive to scaling

Applications and Evaluation Techniques

Applications of Clustering

- Customer segmentation
- Document grouping
- Image segmentation
- Anomaly detection

Applications of PCA

- Data visualization
- Feature reduction
- Compression

Evaluation Techniques for Clustering

Since clustering has no labels, evaluation uses internal metrics.

Silhouette Score

- Measures how well a point fits in its cluster
- Range: -1 to 1
- Higher value = better clustering

WCSS (Within Cluster Sum of Squares)

- Measures compactness of clusters
- Lower value = better clusters

Elbow Method

- Used to find optimal number of clusters
- Choose point where graph bends

External Evaluation (if labels exist)

- Adjusted Rand Index
- Compare predicted clusters with actual labels

Example: Elbow Method

Run K-Means with different K values (2–10).
Plot K vs WCSS.
Choose the point where decrease slows down → optimal K.